

# ***PLOW:***

## ***Probabilistic Logic Over the Well-Founded Semantics***

**Benjamin Grosf and Theresa Swift**

Kyndi, Inc.

<http://benjamingrosof.com>

<http://www3.cs.stonybrook.edu/~tswift>

<https://kyndi.com>



Poster and Demo presentation (20 minutes) at the AAAI 2019 Spring Symposium on  
Combining Machine Learning and Knowledge Engineering (AAAI-MAKE)

March 26, 2019

Stanford, California, USA

Acknowledgements and Thanks: This presentation includes, in part, work done by the authors while at Coherent Knowledge and at Kyndi, Inc. It includes, in part, material originally co-authored by: Michael Kifer, Janine Bloomfield, and Paul Fodor.

Note on Copyright and Permissions:

Slides with Coherent Knowledge logo are, in whole or part, courtesy of Coherent Knowledge; and several other slides as well are, in whole or part, courtesy of Coherent Knowledge.

# Overview of Talk

- Goal: Provide an approach for uncertainty in KRR, to be used in combining logical KRR with ML, that has a better balance of expressiveness and computational scalability.
- Contribution: theory and implementation, as extended form of logic programming
  - Fuzzy (t-norms), in addition to Bayesian
- Simple examples, and brief demo
- How influenced by others' work at the Symposium:
  - Applications / use cases
  - Design patterns for adding KRR to ML

# Motivation and Background

- Probabilistic logic KRR is a fundamental bridge between ML and KE
- Declarative logic programs (LP) is the central KR of IT
  - DBs: Relational DBs (SQL). Knowledge graphs, a.k.a. graph DBs (SPARQL).
  - Ontologies: OWL-RL, RDF-S.
  - Rules: Prolog; RIF; Production rules, Event-Condition-Action rules.
- LP's non-classical logic – invented by/for computer science not math
  - Humble spirit: avoid reasoning-by-cases/disjunction; avoid proof-by-contradiction; stay grounded
  - Well-founded semantics: 3 truth values, benefits for scalability & robustness
- Rulelog – extended LP with high expressiveness + scalability
  - Defeasibility, higher-order syntax, object-oriented (frame) syntax, quantified classical-like formulas, restraint bounded rationality, provenance; poly-time!
  - But lacks kind of quantitative uncertainty needed to reason productively and efficiently using results from a wide variety of ML approaches
- Distribution semantics – extended LP with Bayesian-flavor probability
  - But **lacks good scalability**, due to reintroducing head disjunction

# Why Need Scalability of the Uncertain KRR for Combining ML and KE

- Inner loop of ML
- KB dev edit-test cycle
- Large KGs/KBs

# Presenters' Background

- Kyndi: AI startup combining ML+KRR+NLP; venture-backed
  - Specialized search & question-answering, via advanced knowledge graphs
  - Customers in national intelligence, pharma, other domains
- Benjamin Grosof – Chief Scientist at Kyndi. Previously:
  - Founding CTO/CEO of Coherent Knowledge, AI startup on Rulelog KRR engine
  - Led advanced research portion of Allen Institute for AI's predecessor (Vulcan)
  - MIT Sloan IT professor, DARPA PI, IBM Research projects lead, Accenture exec
  - Co-invented many advances in LP/Rulelog
- Theresa Swift – scientist at Kyndi
  - Also researcher/engineer at US Customs & Border Patrol
  - Lead implementer of XSB
  - Co-founder of Coherent Knowledge
  - Co-invented many advances in LP/Rulelog

# Probabilistic LP – Expressive Extension of LP

- Numerical truth values for atoms (and rules) range on real interval  $[0..1]$
- *head* formula can be:  $\backslash$ or of disjoint atoms/literals whose weights add to 1
  - $\text{friendly}(?x) \sim 0.8 \backslash$ or  $\text{unfriendly}(?x) \sim 0.2 \text{ :- student}(?x).$
- Two major flavors of numerical uncertainty
  1. Bayesian flavor cf. “distribution semantics” [Sato]
    - Superset of Bayesian Networks, expressively
    - General case is computationally intractable, even for function-free
  2. Generalized “triangular norms” (t-norms), a.k.a. fuzzy flavor.
    - Parametrized by choice of the t-norm function  $F$ .
    - $\text{pr}(A \backslash \text{and } B) = F(\text{pr}(A), \text{pr}(B))$ . I.e., “truth-functional” – key to scalability.
    - E.g.,  $F = \min$ . Co-norm for  $\backslash$ or: e.g.,  $\max$ . Same  $F$  is applied to every  $A, B$ .
    - Polynomial time for function-free
    - Generalization:  $F = \text{MinMax}$ , a function on intervals, where the interval is cautious in regard to the potential correlation of  $A$  and  $B$ .

# Bayesian PLP Reasoning: Example

$\text{heads}(\text{Coin}) \sim 0.5 \text{ \or tails}(\text{Coin}) \sim 0.5 \quad \text{: - toss}(\text{Coin}) \text{ \and fair}(\text{Coin}).$

$\text{heads}(\text{Coin}) \sim 0.6 \text{ \or tails}(\text{Coin}) \sim 0.4 \quad \text{: - toss}(\text{Coin}) \text{ \and biased}(\text{Coin}).$

$\text{fair}(\text{Coin}) \sim 0.9 \text{ \or biased}(\text{Coin}) \sim 0.1.$

$\text{toss}(\text{coin}).$

- Conclude:  $\text{heads}(\text{Coin}) \sim 0.51$  .

# T-Norms

- Full Bayesian reasoning is powerful but (computationally) expensive.
- Epistemically, Bayesian probabilities may not be a good way to represent similarity and relevancy distances. We say, more generally: “measures”.
- Hence, T-Norms (Triangular Norms, a generalization of Fuzzy Logic)
  - Godel (i.e., “Min” for conjunction): the measure of  $A \text{ op } B$  expresses perfect correlation (+1) of A and B
  - Lukasiewicz: the measure of  $A \text{ op } B$  expresses negative correlation (-1) of A and B
  - Product: the measure of  $A \text{ op } B$  expresses independence (correlation 0) of A and B
  - “MinMax” (new!): generalizes the measure to an [interval](#) [Lukasiewicz, Godel] expressing an interval of truth, cautious in regard to how much correlation of A and B.



# PLOW System for Probabilistic LP

- The first to implement the generalized t-norm flavor
- Bayesian flavor (a.k.a. distribution semantics), too
- Lattice flavor qualitative uncertainty, too
- Supports  $\neg$  (strong negation)
- Utilizes *undefined* truth value, as do normal LP and Rulelog
- A way to combine deductive reasoning with ML facts and rules
  - E.g., in knowledge graphs
- Implementation extends XSB, and is open source.
  - The PLPs are transformed into normal LP
  - BDDs (Binary Decision Diagrams) are used to collate information from different deduction paths
- In-progress: Aim to integrate tightly with as many Rulelog features as possible. Starting with defeasibility and restraint. Already reusing some of Rulelog's algorithms, theory, implementation!
  - Also in progress: support for running as extension of SWI Prolog, too.

All 3 flavors  
under 1 roof;  
mix-and-match.

# PLOW Uses

- Similarity relations – e.g., two documents may be more or less related
- Vague properties – e.g., a certain person may be more or less “tall”
- Relevancy relations – e.g., a document may be more or less relevant to a query
- Confidence measures – e.g., a document may come from a more or less trusted source
- Lower complexity probability measures – such as “evidential” probabilities

# Strong Negation in PLOW

- Notation:
  - naf**(q) denotes default negation of q. (“not believe” q)
  - neg**(q) denotes strong (a.k.a. explicit) negation of q. (“believe opposite” of q)
- Simple example:
  - $p \sim 0.4$ .
  - $p \sim 0.5$ .
  - $p$  :- undefined.
  - $\text{neg}(p) \sim 0.2$ .

In this case,  $p \sim M$  is

t if  $M \leq 0.5$

u if  $0.5 < M < 0.8$

f if  $0.8 < M \leq 1$

**One can view there as being 3 zones (or bands) of measures having the 3 truth values: a zone for (or where) t, a zone for u, a zone for f.**

# PLOW Paraconsistent/Defeasibility Semantics

- Semantics is an extension of Well-Founded Semantics with Explicit Negation to include quantitative values
  - Uses the coherence principle: strong (i.e., explicit) negation implies default negation.
- Paraconsistent values are mapped to u. This is a kind of defeasible conflict handling.
- Thus, given the assertions:
  - $p \sim 0.6$
  - $\text{neg}(p) \sim 0.6$
- Then conclude that:

$p \sim M$  is:

- t for  $M < 0.4$
- u for  $0.4 \leq M \leq 0.6$
- f for  $0.6 < M \leq 1$

$\text{neg}(p) \sim M$  is:

- t for  $M < 0.4$
- u for  $0.4 \leq M \leq 0.5$
- f for  $0.6 < M \leq 1$

*BRIEF DEMO GOES HERE*

The next few slides are screenshots

*Overall Demo – XSB/PLOW command line; and KB editor (in Emacs)*

[illegible]

## *Example KB (zoomed)*

tae\_madrid~0.75.

aqwe\_madrid~0.6.

neg(tae\_madrid)~0.6.

neg(aqwe\_madrid)~0.75.

# *Start XSB, PLOW; load example KB*

```
root@6e3415616c1b:/# /xsb-src/bin/xsb
[xsb_configuration loaded]
[sysinitrc loaded]
[xsbbrat loaded]

XSB Version 3.8.0 (Three-Buck Chuck) of October 28, 2017
[x86_64-unknown-linux-gnu 64 bits; mode: optimal; engine: slg-wam; scheduling: local]
[Build date: 2019-02-18]

| ?- [tnorm].
[tnorm loaded]
current_tnorm(fuzzy)

yes
| ?- [madrid_simple_v12].
[madrid_simple_v12 loaded]

yes
```



# *Query the example KB, in PLOW*

```
| ?- para_t(aqwe_madrid~PROB).
```

```
PROB = 0.25
```

```
yes
```

```
| ?- para_t(neg(tae_madrid)~PROB).
```

```
PROB = 0.25
```

```
yes
```

```
| ?- para_t(neg(aqwe_madrid)~PROB).
```

```
PROB = 0.4
```

```
yes
```

```
| ?-
```

# Conclusions: Contributions

- Multiple flavors of uncertainty for logic programs, all under one roof
  - Bayesian, i.e., distribution semantics. Both general and restricted.
  - Fuzzy, i.e., t-norms. Highly scalable.
  - Lattice, i.e., qualitative
  - Implementation as extension (package) of XSB, inheriting many good features
- Interval t-norm: MinMax
  - With interpretation of bounds on correlation
- Leverages undefined truth value, and supports unstratified NAF
- Supports strong negation ( $\neg$ ), with basic defeasibility
- Supports well: logical functions, in combination with uncertainty
  - Well-defined: Finite number of finite models, unlike other probabilistic LP approaches. Ensured by restraint + call subsumption (features of XSB).
  - Positioned well to combine with the higher-order syntax (HiLog) feature of Rulelog, useful to represent advanced defeasibility, causality, natural language

# Current and Future Directions

- KRR end:
  - Relate MinMax t-norm to approximation of distribution semantics
  - More on defeasibility including prioritization, argumentation meta-rules
  - Explore and roadmap integration with more/rest of Rulelog features
  - Address idempotence issues for product and Lukasiewicz t-norms. Ideas:
    - Path independence cf. IND. Compilation cf. BDDs/circuits. Human-authored control.
  - Converge syntax with LPAD cf. PITA
- ML end:
  - Pursue relationships to important specific ML techniques. Including for:
    - Distribution semantics. E.g., cplint, Problog, PRISM.
    - Neural network deep learning. E.g., via t-norms.
  - Apply to constructing knowledge graphs from NL + structured info
    - As at Kyndi. E.g., in entity tagging.

# For More Info

- Rulelog detailed tutorial (3 hours) at KR-2018 conference (Oct. 2018):
  - At: <http://benjamingrosof.com/misc-publications/#KR2018RulelogTutorial>
  - It links to:
    - <http://benjamingrosof.com/wp-content/uploads/2018/11/talk-kr2018-rulelog-tutorial-slides-2.pdf>
- Invited talk on: why and how to add KRR to ML (July 2018)
  - At: <http://benjamingrosof.com/misc-publications>

# Thank You

Disclaimer: All brands, logos and products are trademarks or registered trademarks of their respective companies.