

# ***PLOW:***

## ***Probabilistic Logic Over the Well-founded Semantics***

**Benjamin Grosf and Theresa Swift**

Kyndi, Inc.

<http://benjamingrosof.com>

<http://www3.cs.stonybrook.edu/~tswift>

<https://kyndi.com>



Poster and Demo presentation (20 minutes) at the AAAI 2019 Spring Symposium on  
Combining Machine Learning and Knowledge Engineering (AAAI-MAKE)

March 26, 2019

Stanford, California, USA

Acknowledgements and Thanks: This presentation includes, in part, work done by the authors while at Coherent Knowledge and at Kyndi, Inc. It includes, in part, material originally co-authored by: Michael Kifer, Janine Bloomfield, and Paul Fodor.

Note on Copyright and Permissions:

Slides with Coherent Knowledge logo are, in whole or part, courtesy of Coherent Knowledge; and several other slides as well are, in whole or part, courtesy of Coherent Knowledge.

# Overview of Talk

- Goal: Provide an approach for uncertainty in KRR, to be used in combining logical KRR with ML, that has a better balance of expressiveness and computational scalability.
- Contribution: theory and implementation, as extended form of logic programming
  - Fuzzy (t-norms), in addition to Bayesian
- Simple examples, and brief demo
- How influenced by others' work at the Symposium:
  - Applications / use cases
  - Design patterns for adding KRR to ML

# Motivation and Background

- Probabilistic logic KRR is a fundamental bridge between ML and KE
- Declarative logic programs (LP) is the central KR of IT
  - DBs: Relational DBs (SQL). Knowledge graphs, a.k.a. graph DBs (SPARQL).
  - Ontologies: OWL-RL, RDF-S.
  - Rules: Prolog; RIF; Production rules, Event-Condition-Action rules.
- LP's non-classical logic – invented by/for computer science not math
  - Humble spirit: avoid reasoning-by-cases/disjunction; avoid proof-by-contradiction; stay grounded
  - Well-founded semantics: 3 truth values, benefits for scalability & robustness
- Rulelog – extended LP with high expressiveness + scalability
  - Defeasibility, higher-order syntax, object-oriented (frame) syntax, quantified classical-like formulas, restraint bounded rationality, provenance; poly-time!
  - But lacks kind of quantitative uncertainty needed to reason productively and efficiently using results from a wide variety of ML approaches
- Distribution semantics – extended LP with Bayesian-flavor probability
  - But **lacks good scalability**, due to reintroducing head disjunction

# Why Need Scalability of the Uncertain KRR for Combining ML and KE

- Inner loop of ML
- KB dev edit-test cycle
- Large KGs/KBs

# Presenters' Background

- Kyndi: AI startup combining ML+KRR+NLP; venture-backed
  - Specialized search & question-answering, via advanced knowledge graphs
  - Customers in national intelligence, pharma, other domains
- Benjamin Grosof – Chief Scientist at Kyndi. Previously:
  - Founding CTO/CEO of Coherent Knowledge, AI startup on Rulelog KRR engine
  - Led advanced research portion of Allen Institute for AI's predecessor (Vulcan)
  - MIT Sloan IT professor, DARPA PI, IBM Research projects lead, Accenture exec
  - Co-invented many advances in LP/Rulelog
- Theresa Swift – scientist at Kyndi
  - Also researcher/engineer at US Customs & Border Patrol
  - Lead implementer of XSB
  - Co-founder of Coherent Knowledge
  - Co-invented many advances in LP/Rulelog

# Probabilistic LP – Expressive Extension of LP

- Numerical truth values for atoms (and rules) range on real interval  $[0..1]$
- *head* formula can be:  $\backslash$ or of disjoint atoms/literals whose weights add to 1
  - $\text{friendly}(?x) \sim 0.8 \backslash$ or  $\text{unfriendly}(?x) \sim 0.2 \text{ :- student}(?x).$
- Two major flavors of numerical uncertainty
  1. Bayesian flavor cf. “distribution semantics” [Sato]
    - Superset of Bayesian Networks, expressively
    - General case is computationally intractable, even for function-free
  2. Generalized “triangular norms” (t-norms), a.k.a. fuzzy flavor.
    - Parametrized by choice of the t-norm function  $F$ .
    - $\text{pr}(A \backslash \text{and } B) = F(\text{pr}(A), \text{pr}(B))$ . I.e., “truth-functional” – key to scalability.
    - E.g.,  $F = \min$ . Co-norm for  $\backslash$ or: e.g.,  $\max$ . Same  $F$  is applied to every  $A, B$ .
    - Polynomial time for function-free
    - Generalization:  $F = \text{MinMax}$ , a function on intervals, where the interval is cautious in regard to the potential correlation of  $A$  and  $B$ .

# Bayesian PLP Reasoning: Example

$\text{heads}(\text{Coin}) \sim 0.5 \text{ \or tails}(\text{Coin}) \sim 0.5 \quad \text{: - toss}(\text{Coin}) \text{ \and fair}(\text{Coin}).$

$\text{heads}(\text{Coin}) \sim 0.6 \text{ \or tails}(\text{Coin}) \sim 0.4 \quad \text{: - toss}(\text{Coin}) \text{ \and biased}(\text{Coin}).$

$\text{fair}(\text{Coin}) \sim 0.9 \text{ \or biased}(\text{Coin}) \sim 0.1.$

$\text{toss}(\text{coin}).$

- Conclude:  $\text{heads}(\text{Coin}) \sim 0.51 .$

# T-Norms

- Full Bayesian reasoning is powerful but (computationally) expensive.
- Epistemically, Bayesian probabilities may not be a good way to represent similarity and relevancy distances. We say, more generally: “measures”.
- Hence, T-Norms (Triangular Norms, a generalization of Fuzzy Logic)
  - Godel (i.e., “Min” for conjunction): the measure of  $A \text{ op } B$  expresses perfect correlation (+1) of A and B
  - Lukasiewicz: the measure of  $A \text{ op } B$  expresses negative correlation (-1) of A and B
  - Product: the measure of  $A \text{ op } B$  expresses independence (correlation 0) of A and B
  - “MinMax” (new!): generalizes the measure to an [interval](#) [Lukasiewicz, Godel] expressing an interval of truth, cautious in regard to how much correlation of A and B.



# PLOW System for Probabilistic LP

- The first to implement the generalized t-norm flavor
- Bayesian flavor (a.k.a. distribution semantics), too
- Lattice flavor qualitative uncertainty, too
- Supports  $\neg$  (strong negation)
- Utilizes *undefined* truth value, as do normal LP and Rulelog
- A way to combine deductive reasoning with ML facts and rules
  - E.g., in knowledge graphs
- Implementation extends XSB
  - The PLPs are transformed into normal LP
  - BDDs (Binary Decision Diagrams) are used to collate information from different deduction paths
- In-progress: Aim to integrate tightly with as many Rulelog features as possible. Starting with defeasibility and restraint. Already reusing some of Rulelog's algorithms, theory, implementation!

# PLOW Uses

- Similarity relations – e.g., two documents may be more or less related
- Vague properties – e.g., a certain person may be more or less “tall”
- Relevancy relations – e.g., a document may be more or less relevant to a query
- Confidence measures – e.g., a document may come from a more or less trusted source
- Lower complexity probability measures – such as “evidential” probabilities

# Strong Negation in PLOW

- Notation:
  - naf**(q) denotes default negation of q. (“not believe” q)
  - neg**(q) denotes strong (a.k.a. explicit) negation of q. (“believe opposite” of q)
- Simple example:
  - $p \sim 0.4$ .
  - $p \sim 0.5$ .
  - $p$  :- undefined.
  - $\text{neg}(p) \sim 0.2$ .

In this case,  $p \sim M$  is

t if  $M \leq 0.5$

u if  $0.5 < M < 0.8$

f if  $0.8 < M \leq 1$

**One can view there as being 3 zones (or bands) of measures having the 3 truth values: a zone for (or where) t, a zone for u, a zone for f.**

# PLOW Paraconsistent/Defeasibility Semantics

- Semantics is an extension of Well-Founded Semantics with Explicit Negation to include quantitative values
  - Uses the coherence principle: strong (i.e., explicit) negation implies default negation.
- Paraconsistent values are mapped to u. This is a kind of defeasible conflict handling.
- Thus, given the assertions:
  - $p \sim 0.6$
  - $\text{neg}(p) \sim 0.6$
- Then conclude that:

$p \sim M$  is:

- t for  $M < 0.4$
- u for  $0.4 \leq M \leq 0.6$
- f for  $0.6 < M \leq 1$

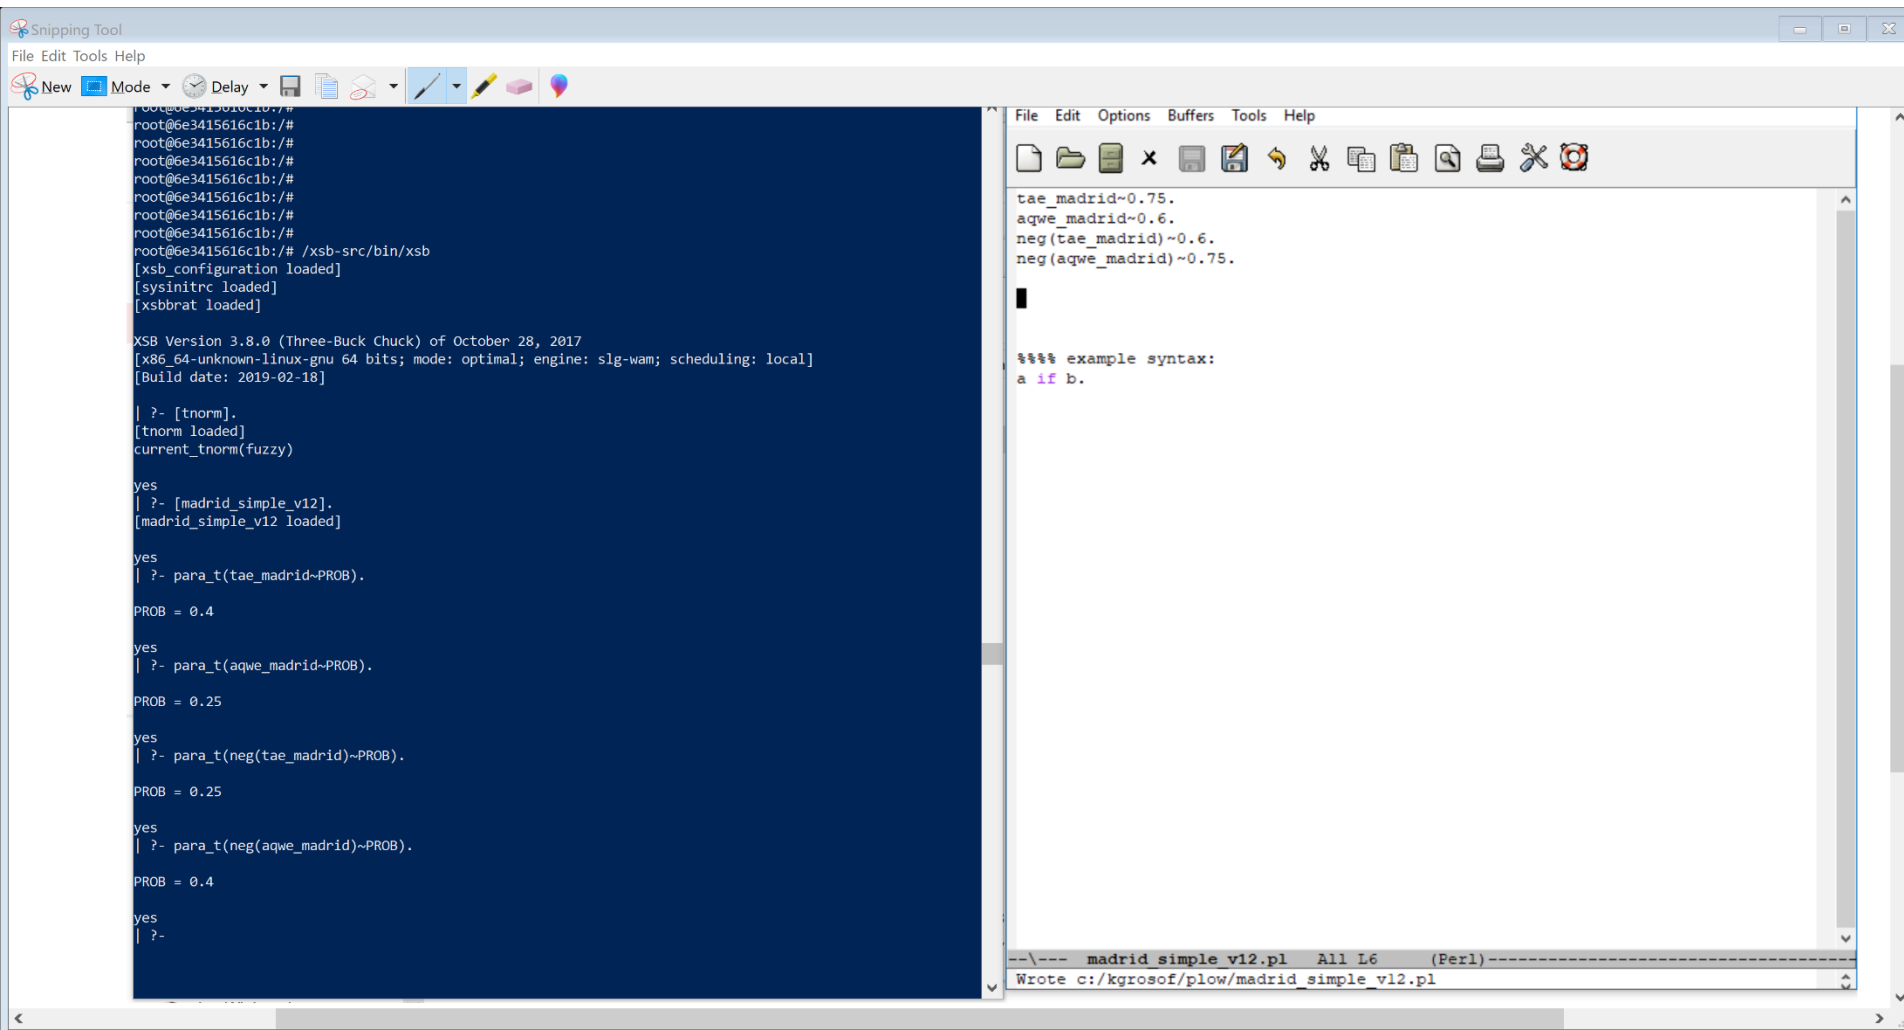
$\text{neg}(p) \sim M$  is:

- t for  $M < 0.4$
- u for  $0.4 \leq M \leq 0.5$
- f for  $0.6 < M \leq 1$

*BRIEF DEMO GOES HERE*

The next few slides are screenshots

# Overall Demo – XSB/PLOW command line; and KB editor (in Emacs)



The screenshot shows a Snipping Tool window with two terminal windows. The left terminal window is running XSB commands, and the right terminal window is running Emacs commands.

```
root@6e3415616c1b:/#
root@6e3415616c1b:/#
root@6e3415616c1b:/#
root@6e3415616c1b:/#
root@6e3415616c1b:/#
root@6e3415616c1b:/#
root@6e3415616c1b:/#
root@6e3415616c1b:/# /xsb-src/bin/xsb
[xsb_configuration loaded]
[sysinitrc loaded]
[xsbbrat loaded]

XSB Version 3.8.0 (Three-Buck Chuck) of October 28, 2017
[x86_64-unknown-linux-gnu 64 bits; mode: optimal; engine: slg-wam; scheduling: local]
[Build date: 2019-02-18]

| ?- [tnorm].
[tnorm loaded]
current_tnorm(fuzzy)

yes
| ?- [madrid_simple_v12].
[madrid_simple_v12 loaded]

yes
| ?- para_t(tae_madrid~PROB).
PROB = 0.4

yes
| ?- para_t(aqwe_madrid~PROB).
PROB = 0.25

yes
| ?- para_t(neg(tae_madrid)~PROB).
PROB = 0.25

yes
| ?- para_t(neg(aqwe_madrid)~PROB).
PROB = 0.4

yes
| ?-
```

```
File Edit Options Buffers Tools Help
tae_madrid~0.75.
aqwe_madrid~0.6.
neg(tae_madrid)~0.6.
neg(aqwe_madrid)~0.75.

**** example syntax:
a if b.

--\--- madrid_simple_v12.pl All L6 (Perl)-----
Wrote c:/kgrosof/pLOW/madrid_simple_v12.pl
```

## *Example KB (zoomed)*

tae\_madrid~0.75.

aqwe\_madrid~0.6.

neg(tae\_madrid)~0.6.

neg(aqwe\_madrid)~0.75.

# *Start XSB, PLOW; load example KB*

```
root@6e3415616c1b:/# /xsb-src/bin/xsb
[xsb_configuration loaded]
[sysinitrc loaded]
[xsbbrat loaded]

XSB Version 3.8.0 (Three-Buck Chuck) of October 28, 2017
[x86_64-unknown-linux-gnu 64 bits; mode: optimal; engine: slg-wam; scheduling: local]
[Build date: 2019-02-18]

| ?- [tnorm].
[tnorm loaded]
current_tnorm(fuzzy)

yes
| ?- [madrid_simple_v12].
[madrid_simple_v12 loaded]

yes
```



# *Query the example KB, in PLOW*

```
| ?- para_t(aqwe_madrid~PROB).
```

```
PROB = 0.25
```

```
yes
```

```
| ?- para_t(neg(tae_madrid)~PROB).
```

```
PROB = 0.25
```

```
yes
```

```
| ?- para_t(neg(aqwe_madrid)~PROB).
```

```
PROB = 0.4
```

```
yes
```

```
| ?-
```

# Conclusions: Contributions

- Multiple flavors of uncertainty for logic programs, all under one roof
  - Bayesian, i.e., distribution semantics. Both general and restricted.
  - Fuzzy, i.e., t-norms. Highly scalable.
  - Lattice, i.e., qualitative
  - Implementation as extension (package) of XSB, inheriting many good features
- Interval t-norm: MinMax
  - With interpretation of bounds on correlation
- Leverages undefined truth value, and supports unstratified NAF
- Supports strong negation ( $\neg$ ), with basic defeasibility
- Supports well: logical functions, in combination with uncertainty
  - Well-defined: Finite number of finite models, unlike other probabilistic LP approaches. Ensured by restraint + call subsumption (features of XSB).
  - Positioned well to combine with the higher-order syntax (HiLog) feature of Rulelog, useful to represent advanced defeasibility, causality, natural language

# Current and Future Directions

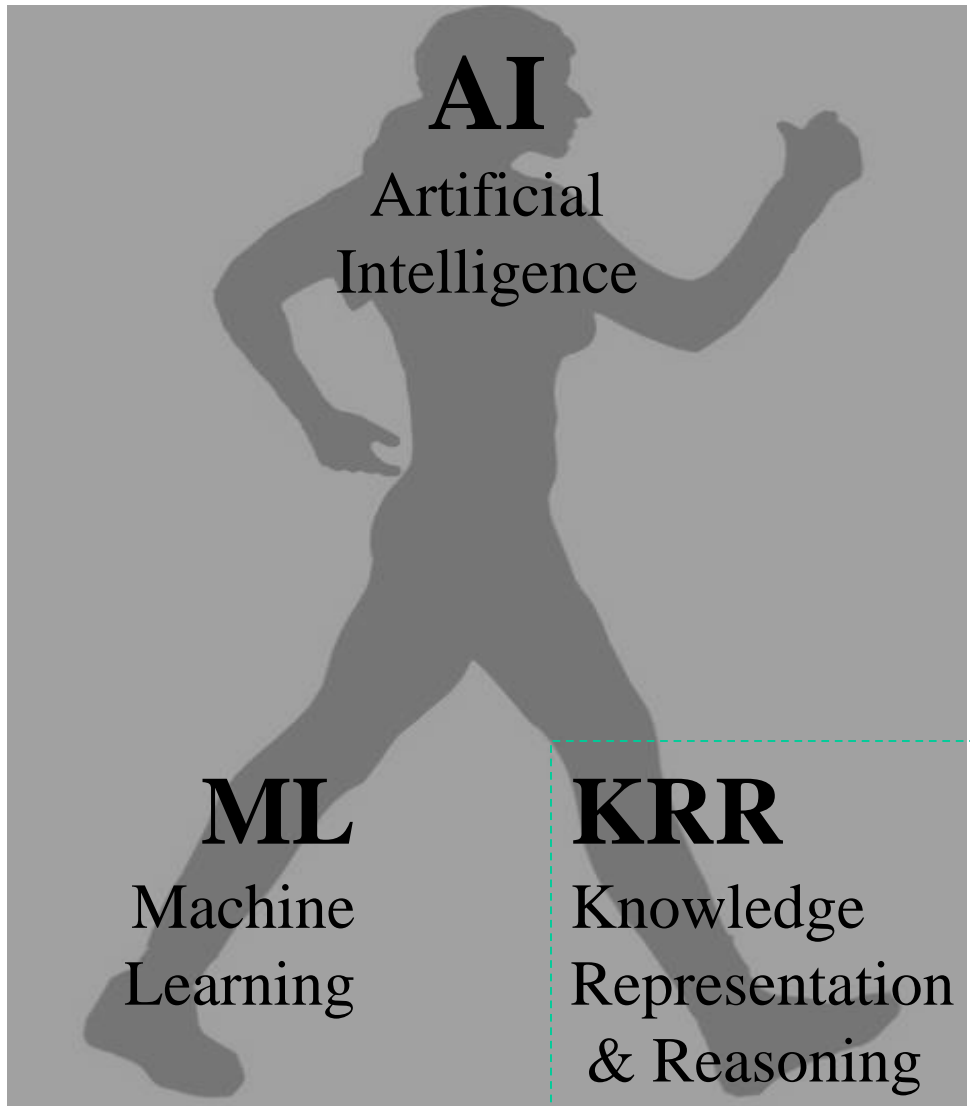
- KRR end:
  - Relate MinMax t-norm to approximation of distribution semantics
  - More on defeasibility including prioritization, argumentation meta-rules
  - Explore and roadmap integration with more/rest of Rulelog features
  - Address idempotence issues for product and Lukasiewicz t-norms. Ideas:
    - Path independence cf. IND. Compilation cf. BDDs/circuits. Human-authored control.
  - Converge syntax with LPAD cf. PITA
- ML end:
  - Pursue relationships to important specific ML techniques. Including for:
    - Distribution semantics. E.g., cplint, Problog, PRISM.
    - Neural network deep learning. E.g., via t-norms.
  - Apply to constructing knowledge graphs from NL + structured info
    - As at Kyndi. E.g., in entity tagging.

# For More Info

- Rulelog detailed tutorial (3 hours) at KR-2018 conference (Oct. 2018):
  - At: <http://benjamingrosof.com/misc-publications/#KR2018RulelogTutorial>
  - It links to:
    - <http://benjamingrosof.com/wp-content/uploads/2018/11/talk-kr2018-rulelog-tutorial-slides-2.pdf>
- Invited talk on: why and how to add KRR to ML (July 2018)
  - At: <http://benjamingrosof.com/misc-publications>

*OPTIONAL SLIDES FOLLOW*

# The Core of AI



*It takes two legs to walk*

# *Logic-based KRR's Roles in AI*

- Complements ML ... in sense of induction from data ... to enable ML in broader sense
- The power of cultural transmission
  - “Evolution’s lesson” (Wolfgang Bibel)
- Accumulate knowledge coherently
- Communicate with humans: expertise, questions
- “Inject” ML results into predictable software

# A Mission for AI Reasoning

- Goal to automate:
  - Powerful reasoning and knowledge that combines tightly:
    1. Multiple features within KRR: uncertainty, defeasibility, ...
    2. KRR with ML. I.e., deductive and inductive inference.
    3. KRR+ML with NLP. For NL understanding and dialogue.
  - Flexible, deep. Scalable. Composable. Evolvable. Explainable.
- *Thereby enable*:
  - Harness humanity's storehouse of textual knowledge
  - Author at low cost: subject matter expertise
  - Star Trek computer interface
  - Unlock \$Trillions of social value in education, healthcare, HCI, customer care, finance & accounting, legal, e-commerce & shopping, science, analytics, workflow & BPM, ...



# Rulelog - Expressive Extension of LP

- Cleanly semantic. Overall: reduces by efficient transformation to normal LP.
- Higher-order syntax (HiLog), reification; rule id's, provenance
  - Higher-order relies on (logical) functions. Elegant transformation.
- Defeasibility: prioritized defaults, exceptions, argumentation
  - Flexible behavior. Efficient approach. Elegant higher-order “argumentation rules”.
  - Both weak negation (`\naf`) and strong negation (`\neg`); `\naf` is outside of `\neg`
- Restraint bounded rationality: guarantee polynomial time
  - Specify *undefined*-ness in various circumstances, e.g., when term size > threshold.
- General classical-like formulas:
  - Head quantifiers; `\exists` treated via skolemization
  - Head `\or`, treated as “omnidirectional” (weak)
- Object-oriented (“frame”) syntax
- External queries; import of most kinds of enterprise info ; thus orchestration
- Probabilistic, to a rudimentary extent, via: higher-order, defeasibility
  - Flexible to “roll your own”: can have tuple of parameters for the probability
    - `pr(formula1)[low->0.91,hi->0.94]. pr(formula2)[mu->0.925,sigma->0.008].`
- Standardization is in process (thru RuleML to W3C and Oasis)

# T-norms (Example of Syntax) in PLOW

```
container_contains_hts(Container,HTS) if  
    cargo_description_hts(Container,HTS),  
    container_xray(Container,HTS).
```

```
cargo_description_hts(c1,7208)~0.2.  
cargo_description_hts(c1,8481)~0.5.
```

```
container_xray(c1,7208)~0.2.  
container_xray(c1,8481)~0.1.
```

# Skipping Restraint – Example in Process Persistence

$\text{occurs}(e1,0) \sim 1.$

$\text{occurs}(E,T) \sim P \text{ :- } T \geq 0, \text{integer}(T), \text{horizon}(\text{Bound}), T \leq \text{Bound},$   
 $T1 \text{ is } T-1, \text{occurs}(E, T1) \sim P1 \text{ } P \text{ is } 0.99 * P1.$

$\text{occurs}(E,T) \sim 0.5 \text{ :- } \text{horizon}(\text{Bound}), T > \text{Bound},$   
 $\text{occurs}(E, \text{Bound}) \sim \_ , T \leq 2 * \text{Bound}.$

$\text{occurs}(E,T) \sim 1 \text{ :- } \text{horizon}(\text{Bound}), T > 2 * \text{Bound},$   
 $\text{occurs}(E, \text{Bound}) \sim \_, \text{undefined}.$

# Reasoning with Uncertainty using Rulelog

- Treat probabilistic statements as a special case of general logical statements and approximate the desired type of reasoning by incorporating numerical uncertainty weights into the general Ergo reasoning facilities. I.e., “roll-your-own” uncertain reasoning. Axiomatize the semantics in Rulelog itself.
  - Leverages highly expressive features of Rulelog including higher-order syntax
  - Examples of reasoning with uncertainty that can be incorporated into rules:
    - Lower and upper bounds on probability value.
      - E.g., *prob\_range(needs\_repair(part32), 0.89, 0.94)*.
    - Confidence level. E.g., *prob\_range\_with\_confidence(needs\_repair(part32), 0.89, 0.94, 0.001)*.
    - Source and provenance info about the statistical or ML method used to derive probability value/range. This can be a basis for numerical uncertainty weights.
      - E.g., *source(prob\_range(needs\_repair(part32), 0.89, 0.93), ML\_episode(myFavoriteMLClassifier, 'Michael Kifer', 'Feb 11, 2017', <http://mycompany.com/dataset41>))*.

# Additional Ways of Reasoning with Uncertainty in ErgoAI

In development:

- “Evidential reasoning”: combines information about probabilities based on different conditions and associated data sets
  - Addresses the probability whether a particular “thing” (i.e, a person, situation, etc.) belongs to a particular class.
  - Examples:
    - The probability that person X has a given disease
    - The probability that a given transaction is risky
    - The probability that a given airplane part will fail given its age, type, and manufacturer
  - This type of reasoning does not require complete or even consistent knowledge about probabilities and is scalable

# Logic Programs: technical overview (I)

- Knowledge base (KB) is a set of rules, each of form:
  - *Head\_formula* :- *Body\_formula*.
  - Intuitively: OK to infer (establish) the head if can infer the body
- Basic normal LP: each rule has form:
  - *atom* :- *literal\_1* \and ... \and *literal\_m*.
  - Plus: atoms are all first-order
- *atom* has form: (predicate(arg\_1,...,arg\_k) ), where each arg\_i is a term
- *literal* has form: (atom) or (\naf atom)
- Weak negation: \naf p – p is not believed (essentially, known to be not provable)
- Strong negation: \neg p – p is believed to be strongly false (i.e., opposite of true)
  - Not permitted in normal LP. But permitted in extensions of normal LP, e.g., in Rulelog.
- Aggregation: setof{?x | *condition*}, where ?x appears in *condition* formula
  - Enabled by \naf. Aggregate operators also include avg, max, min, count.
  - average\_salary(?co,?amt) :-  
    company(?co) \and avg(?amt | employee(?co,?e) \and salary(?e,?amt)).

# Logic Programs: technical overview (II)

- Horn subset: body literals are restricted to be atoms
- Datalog subset: Horn, and function-free
- Full normal LP permits also:
  - in head: `\and`
  - in body, freely nested: `\or`, `\forall`, `\exists`, aggregates, `\and`, `\naf`
  - Integrity constraints via *violation(...)* as a head atom predicate
  - Reduces via transformation to basic normal LP
- Semantics (well-founded) is based on:
  - An alternating least fixed point construction in 3-valued logic
  - Each instantiated atom is assigned to 1 of 3 truth values {t,f,u}:
    - $t = \text{true}$ ;  $f = \text{"false"}$  (cf. `\naf`);  $u = \text{"undefined"}$  (don't-care).
    - *undefined* is useful for paradox and restraint bounded rationality
- Function-free case is polynomial time
- Functions (thru recursion) lead to undecidability

# *LP is the Central Form of Practical Logic*

- *LP is the core KR of structured knowledge management today*
- A non-classical logic invented by computer scientists
- Subsets of LP – important in industry landscape today
  - Relational databases (SQL)    *[Datalog subset of LP]*
  - Knowledge graphs, a.k.a. graph databases (SPARQL)    *[Datalog]*
    - Also: XML databases, object-oriented databases, other semi-structured databases
  - Production rules, Event-Condition-Action rules.    More precisely: their logical subsets.
  - Prolog.    More precisely: its “pure” logical subset.
  - Ontologies, incl. for knowledge graphs: e.g., OWL-RL, RDF-Schema *[Datalog]*
  - *Industry standards for semantic rules*
    - Many RuleML & RIF dialects, e.g., RIF-BLD, RIF-Core, SWRL



# *Some State-of-the-Art Semantic KRR Systems*

- LP: XSB (Stonybrook U., Theresa Swift, David Warren, et al)
  - Full programming language that is Prolog+
- Rulelog: ErgoAI (Coherent Knowledge, free for research), and its open-source subset Flora-2 (originally Stonybrook U.)
  - Full programming language that is Prolog++ and XSB++
- Probabilistic LP: Problog (Luc de Raedt et al); PRISM (T. Sato et al); PLOW (T. Swift, F. Riguzzi, B. Grosz)
- Probabilistic Soft Logic: (Lise Getoor et al, UC Santa Cruz)
- Markov Logic Networks: Alchemy (Pedro Domingos et al, U. Washington)
- Probabilistic Graphical Models generally: See STARAI workshops

# Computation

- XSB has an engine that computes WFS in a scalable manner
  - Applications with up to  $O(10^8)$  facts kept in XSB's space. You can of course scale more using database interfaces (or more RAM!)
- The ErgoAI system uses a highly sophisticated compiler from Rulelog to XSB that supports sophisticated program constructs, including adjoint defeasibility theories, quantifiers, a frame-based syntax and much else
- The XSB engine also supports a lattice-theoretic semantics – all of the quantitative logics mentioned above can be modeled by using different lattices. (Bayesian probabilities requires additional infrastructure such as Binary Decision Diagram libraries.)

*END OF OPTIONAL SLIDES*

# Thank You

Disclaimer: All brands, logos and products are trademarks or registered trademarks of their respective companies.